



Manuel Pérez Martínez

Cloud Platform Architect · AI Platform and Infrastructure

Location Gijón, Spain · remote
Email manugijon@gmail.com
Web manuelpm.com/me · scoutr.gg

Phone (+34) 660 163 565
LinkedIn linkedin.com/in/mperezmartin
GitHub github.com/manupm87

PROFILE

Cloud platform architect and engineer with 11 years in SaaS, fintech and telecom R&D, on Kubernetes since 2015 and Google Cloud in production since 2019. I design platforms, define them as code and run them in production, with cost as a design criterion: cost allocation by business unit, spend reduction and the unit cost of AI workloads. Production on GCP (GKE); I have also built AWS infrastructure.

In 2026 I founded Scoutr.gg, a real-time AI product with paying subscribers, where I measure the inference cost of every game and cap each account's consumption server-side.

CORE SKILLS

FinOps and cost

Cost allocation by label and business unit, billing dashboards, committed-use discounts, right-sizing, observability cost, LLM unit cost and usage quotas

AI and MLOps

Claude API, Amazon Bedrock, Vertex AI, RAG and vector stores, MLflow (Master's), Claude Code, MCP

GCP

GKE, Cloud Run, Compute, VPC networking, load balancing, Cloud Armor, IAM, IAP, VPC Service Controls

AWS

Lambda, API Gateway, CloudFront, S3, DynamoDB, Cognito, IAM, Route 53, ECR, WAF, Bedrock

Kubernetes and platform

Kubernetes, Docker, Helm, Istio, ArgoCD, Argo Rollouts, autoscaling, operators, ephemeral environments

IaC and CI/CD

Terraform, Config Connector, Jsonnet and Tanka, Ansible, GitLab CI, GitHub Actions, Jenkins, GitOps

Reliability and security

SRE and on-call, Datadog, New Relic, Prometheus, Grafana, Sentry, HashiCorp Vault

Programming

Python, TypeScript, Bash

WORK EXPERIENCE

2026 – Present

Founder and Engineer

Scoutr.gg · Remote · scoutr.gg

Real-time AI coaching for League of Legends (computer vision, an LLM and a voice-guided overlay). Built and operated solo; paid and in production since September 2026.

- Built the LLM coaching layer on the Claude API with retrieval-augmented generation, **per-game token accounting**, a shared call budget and server-side limits. Inference cost is measured at €0.05 to €0.07 per coached game.
- Implemented billing with **per-account usage metering**: Paddle subscriptions, tiered entitlements and server-enforced quotas.
- Designed a **hybrid architecture**: a deterministic core ships in the desktop app (Python, OpenCV, Electron), while the AI layer, prompts and keys stay on GCP, on two Cloud Run services with separate trust levels, Firestore and Secret Manager, all in Terraform.
- Built **continuous deployment** on GitHub Actions with Workload Identity Federation (no service account keys) and observability with Sentry, log-based metrics and Grafana. Delivered with Claude Code agents: more than 40 releases since March 2026.

2022 – 2026

Cloud Platform Engineer / Architect

Storyteq · London, United Kingdom (remote from Spain)

- Designed and implemented the company's **cloud platform on GCP and Kubernetes**, and migrated the workloads onto it, in a team of 7: 10+ GKE clusters across EU and US regions, running 20+ microservices per environment, all in Terraform and Config Connector.
- Ran **cloud cost management (FinOps)**: cost allocation by business unit through labels, billing-export dashboards, and savings from committed-use discounts, right-sizing and preemptible nodes.
- Ran **observability on Datadog and New Relic** (dashboards, alerts, APM, log pipelines, SLOs), compared the two vendors on cost, and cut the bill with log exclusion and sampling.
- Rolled out **Anthropic models on Vertex AI** to the engineering team with IAM access control and usage quotas, and provisioned GPU VMs for model training.
- Contributed to the design of the move from a regional (EU) to a **global (EU and US) SaaS offering**, using global external load balancing and Istio as a Gateway API provider.

- Operated the clusters day to day: workload and node-pool autoscaling, resource sizing, and the **production on-call rota** one week in four.
- Designed the **ephemeral environments** that replaced a single shared staging environment (a label on a GitLab merge request deploys the full application in under a minute; used daily by 30+ engineers) and GitOps delivery with ArgoCD, Argo Rollouts and Helm.

2021 – 2022

Site Reliability Engineer (contract)

Ziglu · London, United Kingdom (remote) · FCA-regulated fintech

- Ran the **GKE clusters** behind the live application and the corporate tooling in a team of 4 SREs, with production on-call, and helped design and operate the network (shared VPCs, firewall policies, load balancers, partner VPNs and peering).
- **Hardened security**: least-privilege IAM, zero-trust access with Identity-Aware Proxy, Cloud Armor, and a migration from static secrets to short-lived dynamic secrets in HashiCorp Vault.

2019 – 2021

Cloud Platform Technical Lead

Quantexa · London, United Kingdom

- Joined as Senior Cloud Platform Engineer and was promoted to **lead a team of 5 cloud platform engineers**: platform architecture, the team's workstreams and implementation decisions.
- Owned the **GCP platform** the data analytics product ran on, in production and development: GKE, Compute, networking, security (data-leak protection with VPC Service Controls) and the infrastructure for ETL workloads, automated with Terraform, Ansible and Packer.

2015 – 2019

Senior R&D Engineer

Nokia Bell Labs · Budapest, Hungary

- Built software for **SDN, cloud and orchestration systems** in 5G research on OpenStack, Docker and Kubernetes, including the "5G Connected Cars" demo shown at Mobile World Congress 2016.
- Led a **5-person DevOps team**, setting what to automate and how (Terraform, Ansible, Helm, Jenkins).
- **Lead inventor of a patent** on testing network services (US11171858B2, EP3788744B1) and co-author of a SOFSEM 2018 paper.

2012 – 2015

Software Engineer

DXC Technology (formerly CSC), and freelance · Asturias, Spain

Alfresco, Liferay and SharePoint platforms for regional public administration clients.

PROJECTS

2026

Kyrian World — AI travel planner on AWS

Final team project, Master's at pontia.tech · kyrian-world.com

- **Redesigned the architecture for cost**: replaced a Fargate, load balancer, NAT gateway and RDS design estimated at more than €55 a month with a serverless one that runs for about €5 a month.
- Built the **AWS platform in Terraform** (Lambda behind API Gateway, CloudFront, S3, Cognito, DynamoDB, a WAF, least-privilege IAM) with keyless deployment through OIDC from GitHub Actions.
- Integrated **Amazon Bedrock** for chat and embeddings, with retrieval from an S3 Vectors index, a trace of every AI turn, and a daily token limit per account.

Side project

Genetic Neural Snake

manuelpm.com/gnsnake · Neural networks trained by a genetic algorithm to play Snake.

EDUCATION

2026

Master in Artificial Intelligence, Cloud Computing and DevOps

pontia.tech · completing October 2026

Includes MLOps with MLflow, CI/CD model releases, model serving, Amazon Bedrock and Azure AI agents.

2005 – 2012

MSc Telecommunication Engineering

University of Oviedo · Gijón, Spain · with an Erasmus year at the University of Bologna, Italy

LANGUAGES

Spanish (native) · English (C1, working language since 2015) · Italian (B1).