

Manuel Pérez Martínez

Cloud Platform Engineer / Architect · SRE · AI Platform and Infrastructure**Location** Gijón, Spain · remote**Email** manugijon@gmail.com**Web** manuelpm.com/me · scoutr.gg**Phone** (+34) 660 163 565**LinkedIn** linkedin.com/in/mperezmartin**Nationality** Spanish (EU work rights)

PROFILE

Cloud platform engineer and architect with 11 years in cloud and platform engineering for SaaS, fintech and telecom R&D, on Kubernetes since 2015 and Google Cloud in production since 2019. I design platforms, write IaC and run them in production, with a focus on developer experience and cloud cost. My production clusters ran on GKE, and I have also built AWS infrastructure.

I also build AI products end to end. In 2026 I founded Scoutr.gg, a real-time AI coaching product with paying subscribers, and I deliver it with coding agents: written specs, automated review, and CI/CD on infrastructure I designed.

Interested in senior, principal or architect roles in platform engineering, SRE, and AI platform and infrastructure, remote from Spain.

CORE SKILLS

GCP	GKE, Cloud Run, Compute, VPC networking, load balancing, Cloud Armor, IAM, IAP, VPC Service Controls
AWS	Lambda, API Gateway, CloudFront, S3, DynamoDB, Cognito, IAM, Route 53, ECR, WAF, Bedrock
Kubernetes and platform	Kubernetes, Helm, Istio, ArgoCD, Argo Rollouts, autoscaling, operators, ephemeral environments
IaC and CI/CD	Terraform, Config Connector, Jsonnet and Tanka, Ansible, GitLab CI, GitHub Actions, Jenkins, GitOps
Reliability and security	SRE and on-call, Datadog, New Relic, Prometheus, Grafana, Sentry, HashiCorp Vault, zero trust
FinOps and cost	Cost allocation, billing dashboards, committed-use discounts, right-sizing, LLM cost limits
AI platform and agents	Claude API, Amazon Bedrock, Vertex AI, RAG and vector stores, Claude Code agents, MCP
Languages	Python, TypeScript, Bash

WORK EXPERIENCE

2022 – 2026

Cloud Platform Engineer / Architect

Storyteq · London, United Kingdom (remote from Spain)

- Designed and implemented the company's **cloud platform on GCP and Kubernetes**, and migrated the workloads onto it, in a platform team of 7: 10+ GKE clusters across EU and US regions, running 20+ microservices per environment, all defined in Terraform and Config Connector.
- Contributed to the design of the move from a regional (EU) to a **global (EU and US) SaaS offering**, using global external load balancing, and leveraging Istio as a Gateway API provider.
- Designed the **ephemeral environments** that replaced a single shared staging environment: a label on a GitLab merge request deploys the full application in under a minute, or instantly from a pre-warmed pool. Built on ArgoCD ApplicationSets and Helm, and used daily by 30+ engineers.
- Published golden paths and mentored teams on them, migrated legacy manifests to reusable Helm charts shared across business units, and built **GitOps delivery** with ArgoCD and progressive rollouts with Argo Rollouts.
- Ran **cloud cost management (FinOps)**: cost allocation by business unit through labels, billing-export dashboards, and savings from committed-use discounts, right-sizing and preemptible nodes.
- Operated the clusters day to day: workload and node-pool autoscaling, resource sizing, and the **production on-call rota** one week in four.
- Ran **observability on Datadog and New Relic** (dashboards, alerts, APM, log pipelines, SLOs), compared the two vendors on cost, and cut the bill with log exclusion and sampling.
- Rolled out **Anthropic models on Vertex AI** to the engineering team with IAM access control and usage quotas, after being among the first to adopt AI coding agents, and provisioned GPU VMs for model training.

2026 – Present

Founder and Engineer

Scoutr.gg · Remote · scoutr.gg

Real-time AI coaching companion for League of Legends: computer vision reads the live game, an LLM reasons about it, and a voice-guided overlay coaches the player mid-match. Built, launched and operated solo.

- **Launched publicly in September 2026**, six months after the first commit, with paid subscription plans, currently at 7% conversion.
- Built the **LLM coaching layer** on the Claude API with retrieval-augmented generation, per-game token accounting, a shared call budget and server-side limits. Inference cost is measured at €0.05 to €0.07 per coached game.
- Designed a **hybrid architecture**: a deterministic core ships in the Windows desktop app (Python, OpenCV, Electron, TypeScript), while the AI layer, prompts and keys stay server-side on GCP behind per-user authentication.
- Run the backend on **two Cloud Run services** with separate trust levels, Firestore and Secret Manager, all in Terraform. The public service's account has no IAM roles, so a compromise cannot read data or secrets through its identity.
- Built path-filtered **continuous deployment** on GitHub Actions with Workload Identity Federation (no service account keys) and self-hosted runners, and observability with Sentry, log-based metrics and Grafana.
- Implemented **billing**: Paddle subscriptions and webhooks, tiered entitlements, usage metering and server-enforced quotas.
- Run a **spec-driven delivery workflow** on Claude Code: every change is gated by a GitHub issue, a written spec and a pull request I approve; an implementing agent writes it and parallel agents review it. More than 40 desktop releases since March 2026.

2021 – 2022

Site Reliability Engineer (contract)

Ziglu · London, United Kingdom (remote) · FCA-regulated fintech

- Ran the **GKE clusters** behind the live application and the corporate tooling in a team of 4 SREs, with production on-call, plus Cloud Run and Compute workloads (Envoy proxies, VyOS gateways, Windows servers).
- Helped design and operate the network (shared VPCs, firewall policies, load balancers, partner VPNs and peering) and **hardened security**: least-privilege IAM, zero-trust access with Identity-Aware Proxy, and Cloud Armor.
- Migrated static secrets from GCP Secret Manager to short-lived dynamic secrets in **HashiCorp Vault**.
- Migrated **Helm charts to Jsonnet**, kept cluster state in sync through Tanka and ArgoCD, and ran Argo Workflows pipelines with Prometheus and Grafana monitoring.

2019 – 2021

Cloud Platform Technical Lead

Quantexa · London, United Kingdom

- Joined as Senior Cloud Platform Engineer and was promoted to **lead a team of 5 cloud platform engineers**: architected new platform features, helped define the team's workstreams, decided on implementation approaches and acted as Scrum Master.
- Owned the **GCP platform** the data analytics product ran on, in production and for the engineering teams' development environments: GKE, Compute, Cloud Functions, networking, security (including data-leak protection with VPC Service Controls) and the infrastructure for ETL workloads, with production on-call one week in four.
- Automated provisioning and configuration with **Terraform, Ansible and Packer**, delivery with Jenkins, and monitoring with Prometheus and Grafana.

2015 – 2019

Senior R&D Engineer

Nokia Bell Labs · Budapest, Hungary

- Built software for **SDN, cloud and orchestration systems** in 5G research on OpenStack, Docker and Kubernetes, including the "5G Connected Cars" demo shown at Mobile World Congress 2016.
- Led a **5-person DevOps team**, setting what to automate and how (Terraform, Ansible, Helm, Jenkins).
- **Lead inventor of a patent** on testing network services, granted in the US and Europe, and co-author of a SOFSEM 2018 paper.

2012 – 2015

Software Engineer

DXC Technology (formerly CSC), and freelance · Asturias, Spain

- Developed and administered Alfresco, Liferay and SharePoint platforms, and managed technical projects for regional public administration clients.

PROJECTS

2026

Kyrian World — AI travel planner on AWS

Final team project, Master's at pontia.tech · kyrian-world.com

Main contributor to a web application that streams day-by-day trip itineraries from an LLM, grounded in a city knowledge base, and author of most of its infrastructure.

- Built the **AWS platform in Terraform**: container-image Lambda functions behind API Gateway with response streaming, CloudFront and S3 for the frontend, Cognito sign-in, DynamoDB, Route 53, ECR, a WAF, and one least-privilege IAM role per function.
- **Redesigned the architecture for cost**: replaced a Fargate, load balancer, NAT gateway and RDS design estimated at more than €55 a month with a serverless one that runs for about €5 a month, inside a €30 monthly budget.
- Integrated **Amazon Bedrock** for chat and embeddings, with retrieval from an S3 Vectors index, a trace of every AI turn, and a daily token limit per account.
- Set up **keyless delivery**: GitHub Actions deploys through OIDC, engineers sign in through IAM Identity Center, and no access keys exist.

Side project

Genetic Neural Snake

manuelpm.com/gnsnake

Neural networks trained by a genetic algorithm to play Snake.

EDUCATION

2026

Master in Artificial Intelligence, Cloud Computing and DevOps

pontia.tech · completing October 2026

Includes MLOps pipelines with MLflow and CI/CD model releases, model serving, Amazon Bedrock and Azure AI agent services.

2005 – 2012

MSc Telecommunication Engineering

University of Oviedo · Gijón, Spain · with an Erasmus year at the University of Bologna, Italy

ADDITIONAL

Patent

WO2019210958A1, "Method for testing network services", lead inventor, Nokia Solutions and Networks. Granted as US11171858B2 and EP3788744B1.

Publications and talks

"Multivendor Deployment Integration for Future Mobile Networks", SOFSEM 2018 · "DevOps: A Bell Labs perspective", DevOps Budapest meetup.

Languages

Spanish (native) · English (C1) · Italian (B1), CEFR levels.